




Concurrency in Rust

Alex Crichton

40 Years of Microprocessor Trend Data



Original data up to the year 2010 collected and plotted by M. Horowitz, F. Labonte, O. Shacham, K. Olukotun, L. Hammond, and C. Batten
New plot and data collected for 2010-2015 by K. Rupp

Bug 631527

Parallelize selector matching

[Get help with this page](#)

NEW Assigned to [dzbarsky](#)

▼ **Status** (NEW bug with no priority)


Product: ► Core

Component: ► CSS Parsing and Computation

Status: NEW

Reported: 7 years ago

Reported: 7 years ago



Must be
this tall to
write multi-
threaded
code.



Fearless Concurrency with Rust

Apr 10, 2015 • Aaron Turon

The Rust project was initiated to solve two thorny problems:

- How do you do safe systems programming?
- How do you make concurrency painless?

Initially these problems seemed orthogonal, but to our amazement, the solution turned out to be identical: **the same tools that make Rust safe also help you tackle concurrency head-on.**

Memory safety bugs and concurrency bugs often come down to code accessing data when it shouldn't. Rust's secret weapon is *ownership*, a discipline for access control that systems programmers try to follow but that Rust's compiler checks statically for you.

What Rust has to offer

Strong safety guarantees...

No seg-faults, no data-races, expressive type system.

...without compromising on performance.

No garbage collector, no runtime.

Goal:

Confident, productive systems programming

Concurrency?

Rust?

Libraries

Futures

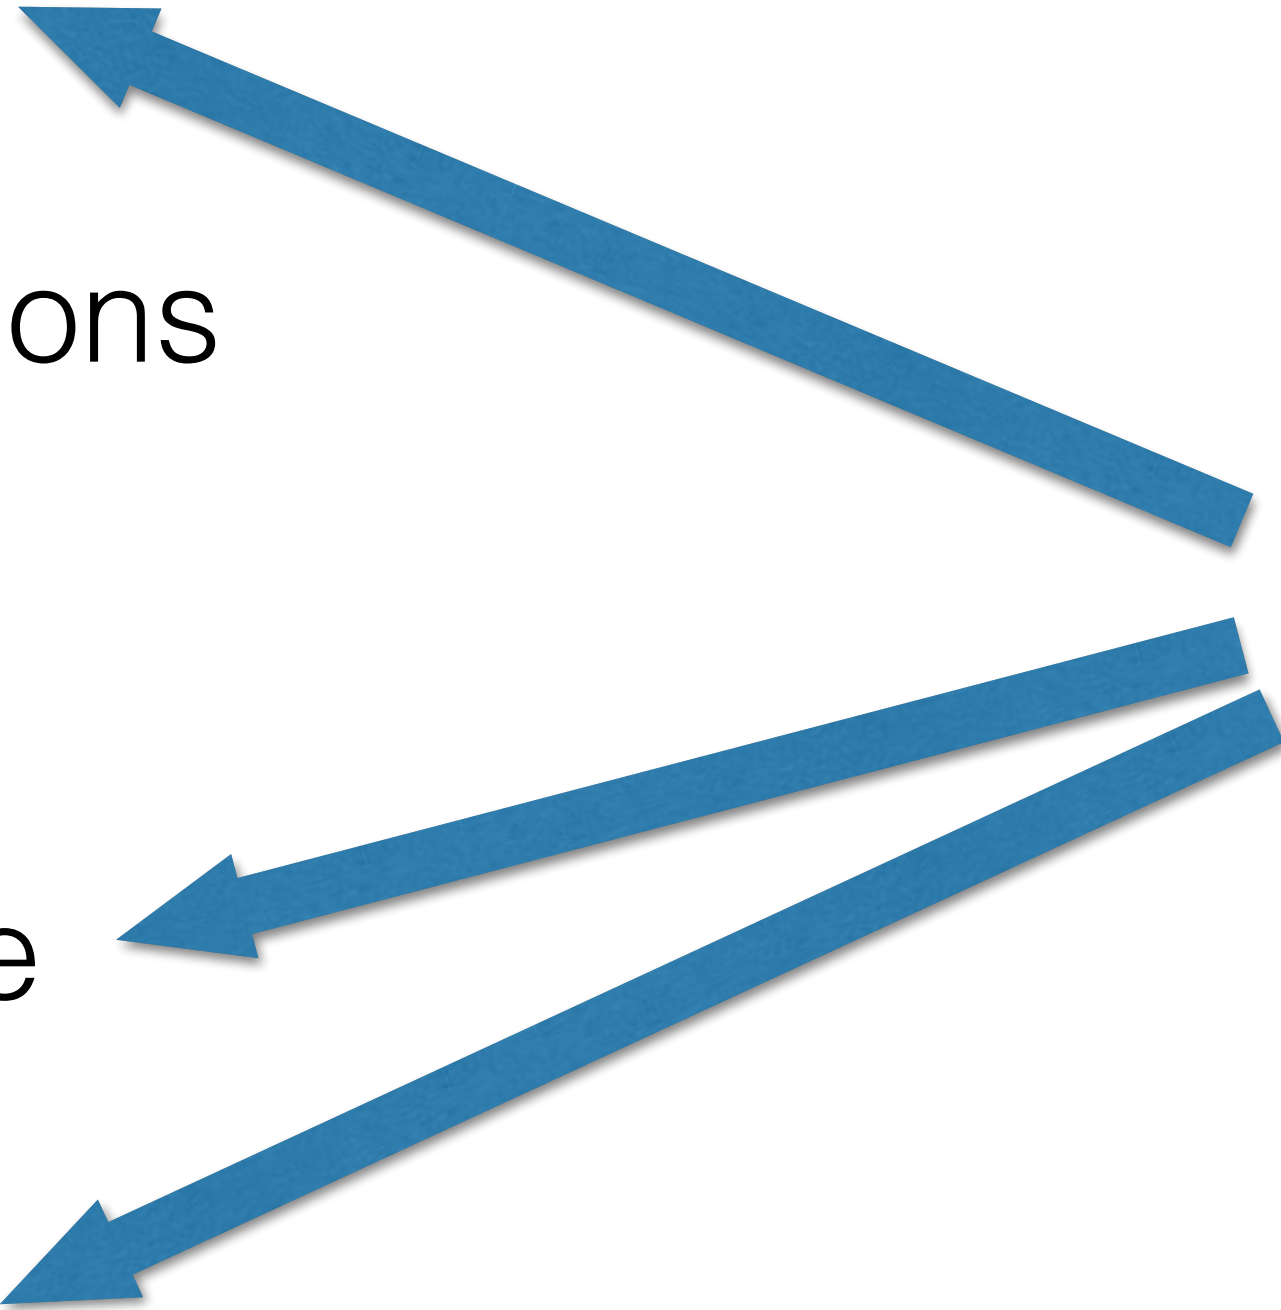
What's concurrency?

In computer science, concurrency is a property of systems in which several computations are executing simultaneously, and potentially interacting with each other.

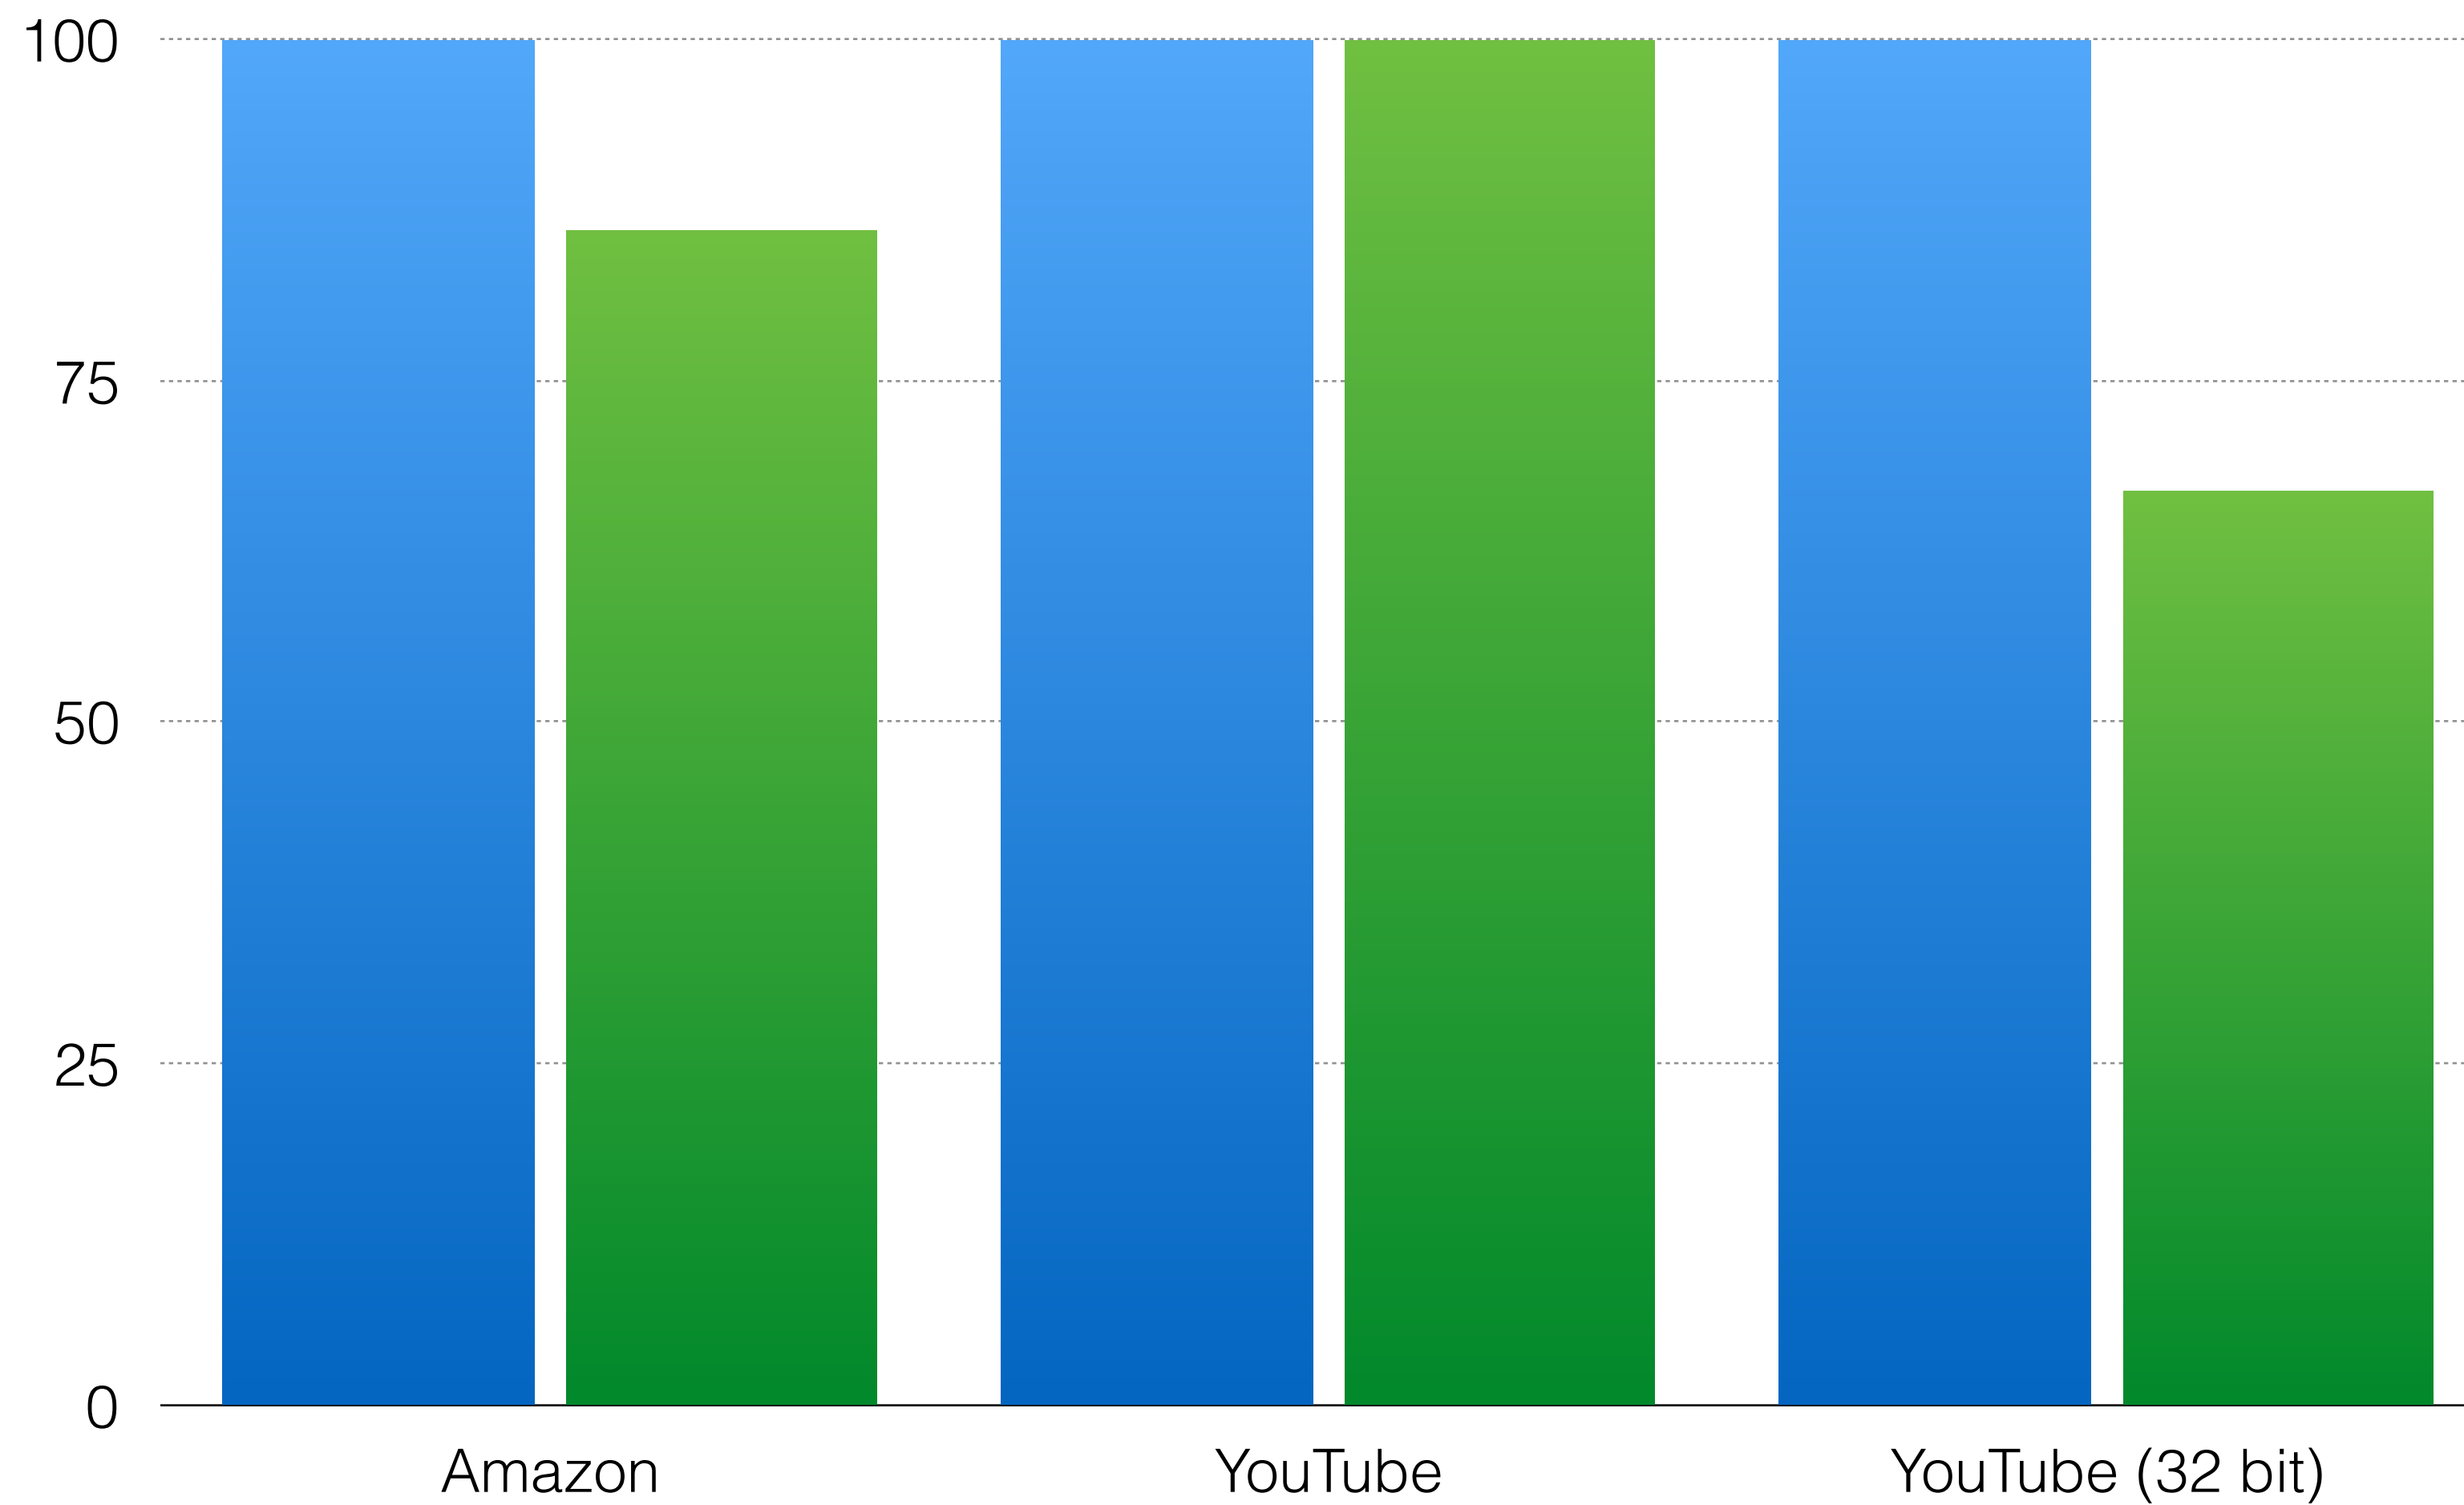
Getting our feet wet

```
// What does this print?  
int main() {  
    int pid = fork();  
    printf("%d\n", pid);  
}
```

Concurrency is hard!

- Data Races
 - Race Conditions
 - Deadlocks
 - Use after free
 - Double free
- Exploitable!
- 
- A diagram consisting of three blue arrows pointing from the word 'Exploitable!' on the right towards a list of concurrency issues on the left. The arrows point to 'Data Races', 'Use after free', and 'Double free'.

Concurrency is nice!



Concurrency?

Rust?

Libraries

Futures

Zero-cost abstractions

+

Memory safety & data-race freedom







=

Confident, productive systems programming

What's **safety**?

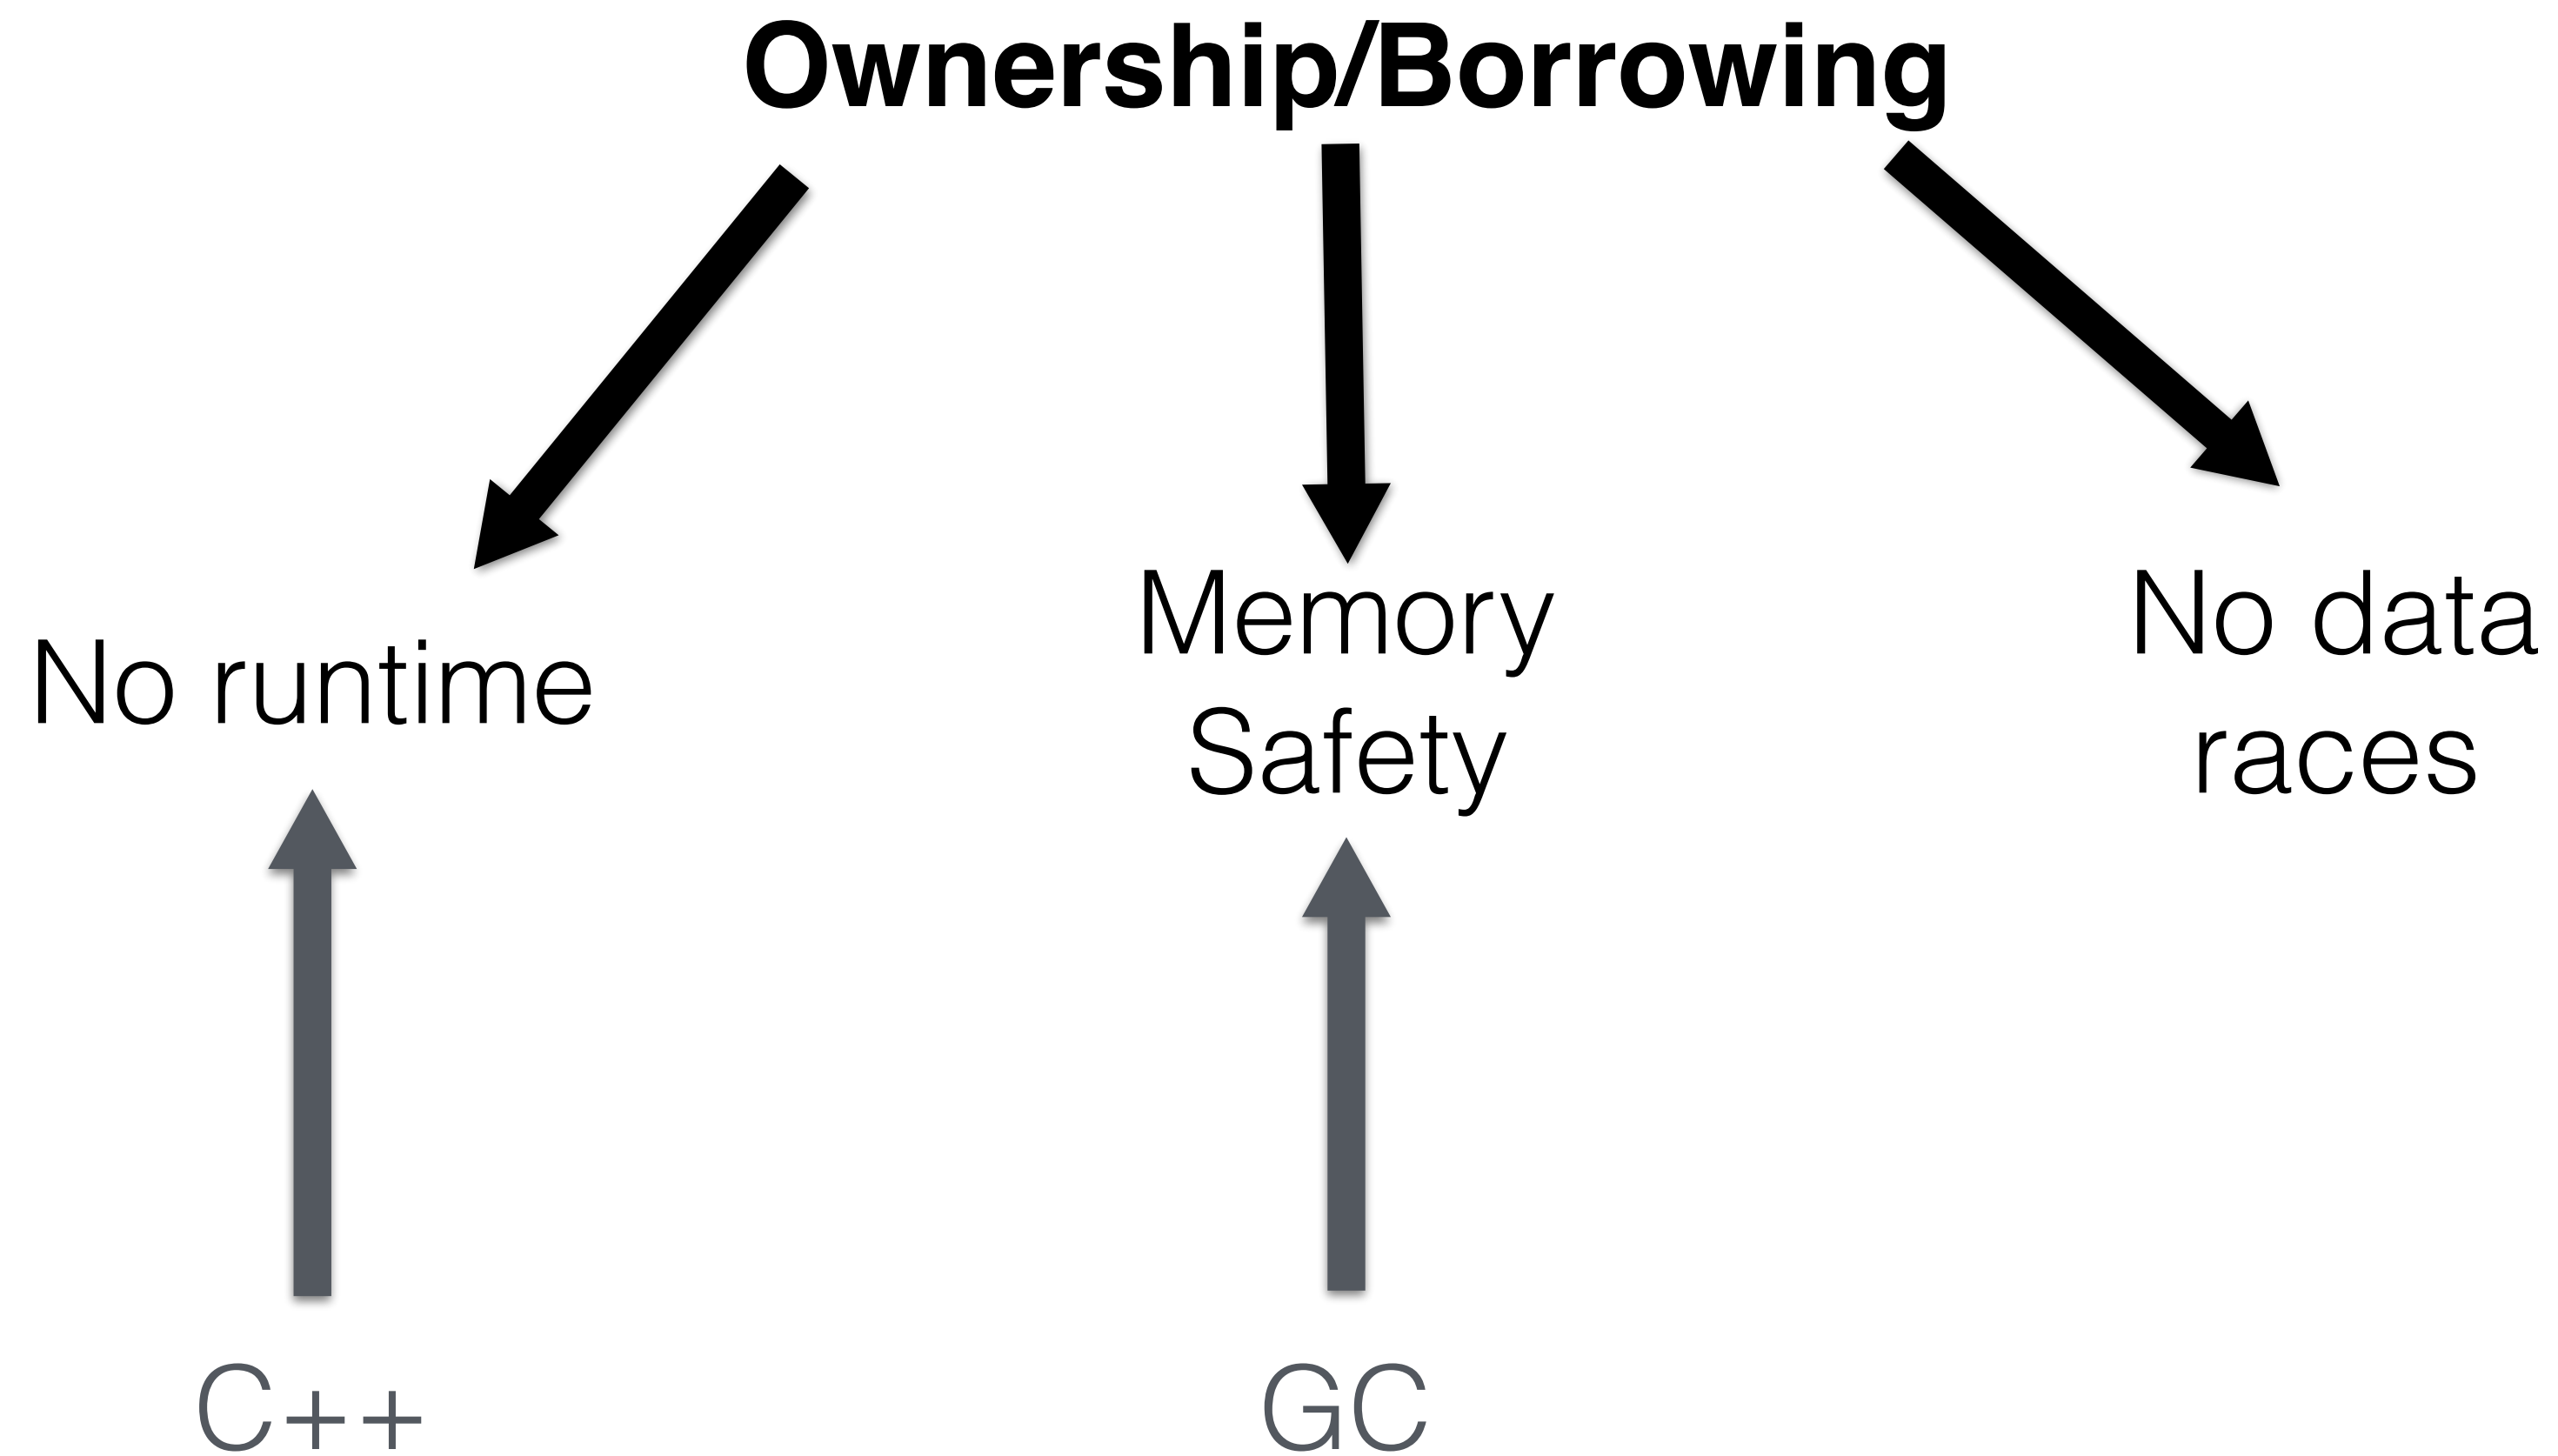
```
void example() {  
    vector<string> vector;  
    // ...  
    auto& elem = vector[0];  
    vector.push_back(some_string);  
    cout << elem;  
}
```

Mutation



Aliased pointers
Dangling pointer!


Rust's Solution



Ownership

```
fn main() {  
  let mut v = Vec::new();  
  v.push(1);  
  v.push(2);  
  take(v);  
  // ...  
}
```

```
fn take(v: Vec<i32>) {  
    // ...  
}
```



Ownership

```
fn main() {  
    let mut v = Vec::new();  
    v.push(1);  
    v.push(2);  
    take(v);  
    // ...  
}  
  
fn take(v: Vec<i32>) {  
    // ...  
}
```


Ownership

```
fn main() {  
    let mut v = Vec::new();  
    v.push(1);  
    v.push(2);  
    take(v);  
    v.push(3);  
}  
  
fn take(v: Vec<i32>) {  
    // ...  
}
```

Ownership

```
fn main() {  
    let mut v = Vec::new();  
    v.push(1);  
    v.push(2);  
    take(v);  
    v.push(3);  
}  
  
fn take(v: Vec<i32>) {  
    // ...  
}
```




error: use of moved value `v`

Borrowing

```
fn main() {  
    let mut v = Vec::new();  
    push(&mut v);  
    read(&v);  
    // ...  
}
```

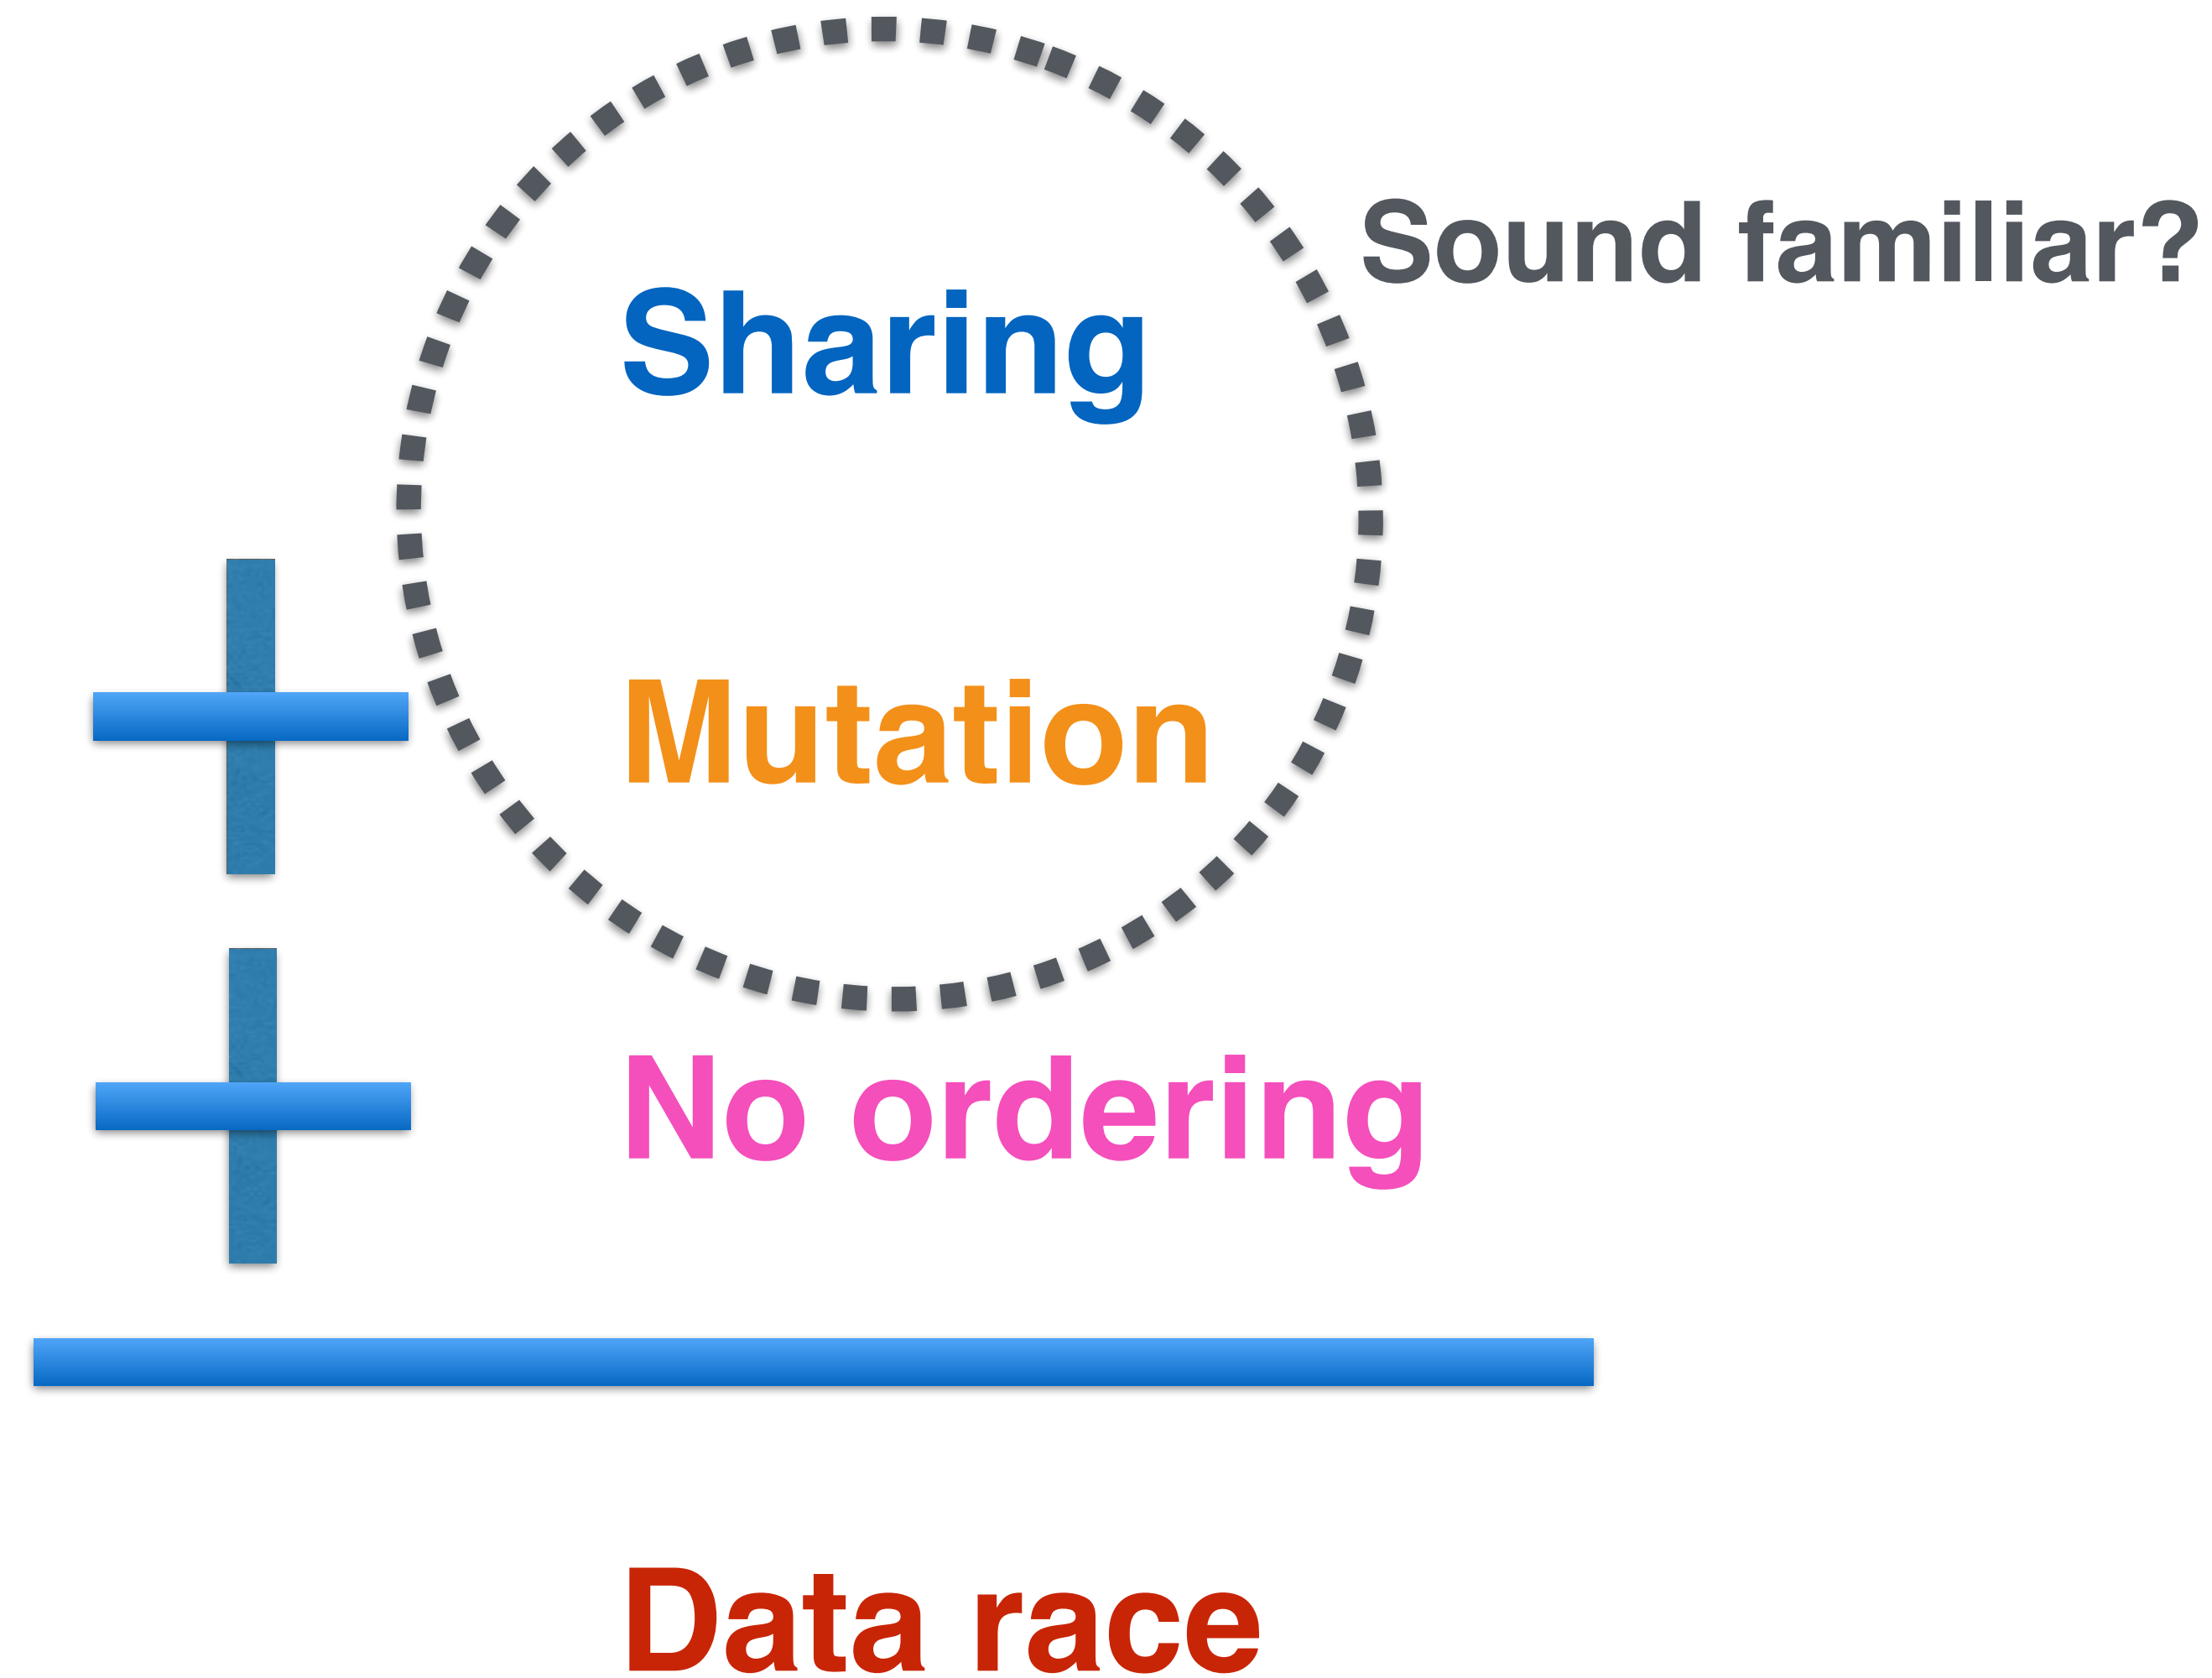
```
fn read(v: &mut Vec<i32>) {  
    // push(1);  
}
```



Safety in Rust

- Rust *statically* prevents aliasing + mutation
- Ownership prevents double-free
- Borrowing prevents use-after-free
- Overall, no segfaults!

Data races



Concurrency?

Rust?

Libraries

Futures

Library-based concurrency

Originally: Rust had message passing built into the language.

Now: library-based, multi-paradigm.

Libraries leverage **ownership and borrowing** to avoid data races.

std::thread

```
let loc = thread::spawn(|| {  
    "world"  
});  
println!("Hello, {}!",  
        loc.join().unwrap());
```

std::thread

```
let mut dst = Vec::new();  
thread::spawn(move || {  
    dst.push(3);  
});
```

```
dst.push(4);
```

error: use after move



std::thread


```
let mut dst = Vec::new();  
thread::spawn(|| {  
    dst.push(3);  
});  
dst.push(4);
```



error: value doesn't live long enough

std::thread


```
let v = Rc::new(vec![1, 2, 3]);  
let v2 = v.clone();  
thread::spawn(move || {  
    println!("{}", v.len());  
});  
another_function(v2.clone());
```



error: Rc<T> can't be sent across threads

std::sync::Arc

```
let v = Arc::new(vec![1, 2]);  
let v2 = v.clone();  
thread::spawn(move || {  
    println!("{}", v.len());  
});  
another_function(&v2);
```



std::sync::Arc

```
let v = Arc::new(vec![1, 2]);  
let v2 = v.clone();  
thread::spawn(move || {  
    v.push(3);  
});  
another_function(&v2);
```




error: cannot mutate through shared reference

std::sync::Mutex

```
fn sync_inc(counter: &Mutex<i32>) {  
    let mut data: Guard<i32> = counter.lock();  
    *data += 1;  
}
```

counter

data




std::sync::atomic::*

```
let number = AtomicUsize::new(10);  
let prev = number.fetch_add(1, SeqCst);  
assert_eq!(prev, 10);  
let prev = number.swap(2, SeqCst);  
assert_eq!(prev, 11);  
assert_eq!(number.load(SeqCst), 2);
```

std::sync::mpsc

```
let (tx, rx) = mpsc::channel();  
let tx2 = tx.clone();  
thread::spawn(move || tx.send(5));  
thread::spawn(move || tx2.send(4));
```

```
// Prints 4 and 5 in an unspecified order  
println!("{:?}", rx.recv());  
println!("{:?}", rx.recv());
```



rayon

```
fn sum_of_squares(input: &[i32]) -> i32 {  
    input.iter()  
        .map(|&i| i * i)  
        .sum()  
}
```

rayon

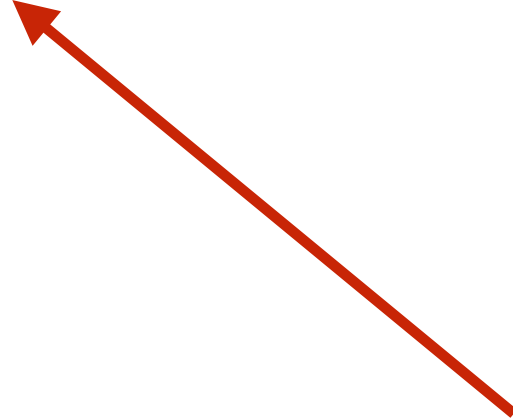
```
use rayon::prelude::*;
```

```
fn sum_of_squares(input: &[i32]) -> i32 {  
    input.par_iter()  
        .map(|&i| i * i)  
        .sum()  
}
```


rayon

```
use rayon::prelude::*;
```

```
fn sum_of_squares(input: &[i32]) -> i32 {  
    let mut cnt = 0;  
    input.par_iter()  
        .map(|&i| {  
            cnt += 1;  
            i * i  
        })  
        .sum()  
}
```



error: `cnt` cannot be shared concurrently

crossbeam

- Epoch-based memory reclamation
- Easy translation of algorithms that require GC
- Work stealing deque
- MPMC queues

100% Safe

- Everything you just saw is foolproof
- No segfaults
- No data races
- No double frees...

Concurrency?


Rust?

Libraries

Futures

Concurrent I/O

- Rust-the-language fuels shared-memory parallelism
- Rust-the-ecosystem fuels other forms of concurrency
- **Futures** are the core foundation for async I/O in Rust



What is async I/O?

```
// blocking  
let amt = socket.read(buffer);  
assert_eq!(amt, 4);
```

```
// async  
let amt = socket.read(buffer);  
assert_eq!(amt, EWOULDBLOCK);
```


What is async I/O?

- No I/O ever blocks
- Application receives bulk lists of events
- User responsible for dispatching events

Async I/O can be hard

User: I'd like the contents of
<https://www.rust-lang.org>, please!

Kernel: file descriptor 5 is ready for reading

Futures

- Sentinel for a value “being computed”
- Internally captures state to produce a value
- Allows composing different kinds of computations


Futures and async I/O

User: I'd like the contents of
<https://www.rust-lang.org>, please!

Library: OK, here's a `Future<Vec<u8>>`

Futures without I/O

```
let result = pool.spawn(|| {  
    fibonacci(100)  
});
```



```
println!("calculating..");  
get_some_coffee();  
run_an_errand();
```

```
let result = result.wait();  
println!("fib(100) = {}", result);
```

Rust's solution: Tokio


cross-platform async I/O

Tokio is the fusion of `mio` and `futures`

you can build futures with a TCP socket

Rust's solution: Tokio

- Futures powered by TCP/UDP/...
- Works on all major platforms
- Implementation of an *event loop*



Futures today

- Traits for one value (Future), many values (Stream), or pushing values (Sink)
- “oneshot channels” - computing an in-memory value asynchronously
- “mpsc channels” - like `std::sync::mpsc`
- Integration with crates such as rayon

Futures today

```
#[async]
fn fetch(client: hyper::Client, url: String)
    -> Result<Vec<u8>, hyper::Error>
{
    let response = await!(client.get(&url));
    if !response.status().is_ok() {
        return Err(...)
    }
    let mut body = Vec::new();
    #[async]
    for chunk in response.body() {
        body.extend(chunk);
    }
    Ok(body)
}
```

Tokio today

- Crates such as `tokio-core`, `tokio-proto`, `tokio-service`
- Considering a `tokio` crate ([tokio-rs/tokio-rfcs#2](https://tokio.rs/tokio-rfcs/#2))
- Supports TCP, UDP, Unix sockets, Named pipes, processes, signals, HTTP, HTTP/2 (coming soon!), websockets, ...
- Deployed to production in a number of companies

Implementing Futures

```
struct Future<T> {  
    // ...  
}
```

What if I have a more efficient implementation?

```
trait Future {  
    type Item;  
  
    // ...  
}
```

```
trait Future {  
    type Item;  
  
    // ...  
}
```

“A future is a sentinel for something being computed”

```
trait Future {  
    type Item;  
  
    fn schedule<F>(self, F)  
        where F: FnOnce(T) ;  
}
```

Can't do virtual dispatch!

```
fn compute(&self, key: u32)
    -> impl Future<Item=u32>
{
    if let Some(v) = self.cache.get(&key) {
        future::ok(v)
    } else {
        self.compute_slow(key)
    }
}
```

Different types!

```
fn compute(&self, key: u32)
    -> impl Future<Item=u32>
{
    if let Some(v) = self.cache.get(&key) {
        future::ok(v)
    } else {
        self.compute_slow(key)
    }
}
```



```

fn compute (&self, key: u32)
    -> impl Future<Item=u32>
{
    if let Some (v) = self.cache.get (&key) {
        Either::A (future::ok (v) )
    } else {
        Either::B (self.compute_slow (key) )
    }
}

```

Assuming:

```

impl<A, B> Future for Either<A, B>
    where A: Future, B: Future

```

```
fn compute (&self, key: u32)
    -> impl Future<Item=u32>
{
    if let Some (v) = self.cache.get (&key) {
        Either::A (future::ok (v) )
    } else if ... {
        Either::B (self.compute_slow (key) )
    } else {
        // ...
    }
}
```

Either::A** ?**

```

fn compute (&self, key: u32)
    -> Box<Future<Item=u32>>
{
    if let Some (v) = self.cache.get (&key) {
        Box::new(future::ok(v))
    } else if ... {
        Box::new(self.compute_slow(key))
    } else {
        // ...
    }
}

```

Virtual dispatch!

```
trait Future {  
    type Item;  
  
    fn schedule<F>(self, F)  
        where F: FnOnce(T) ;  
}
```

Can't do virtual dispatch!

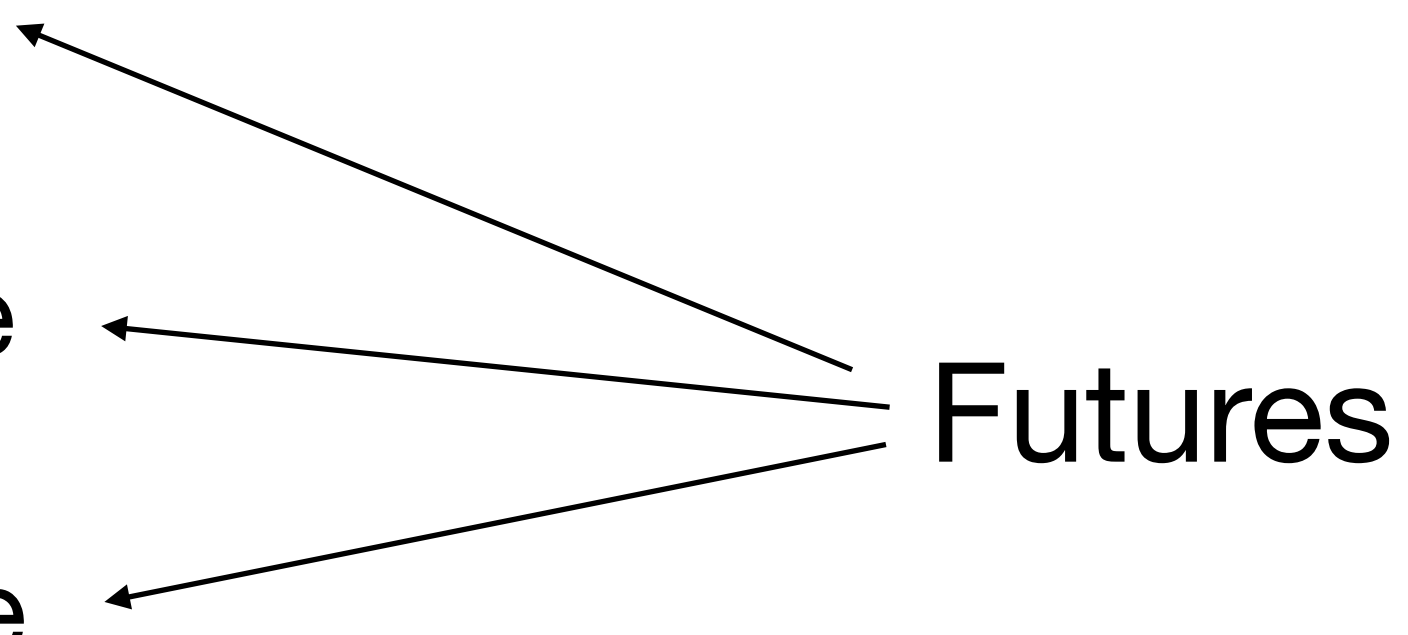
```
trait Future {  
    type Item;  
  
    fn schedule (&mut self,  
                f: Box<FnOnce (T) >) ;  
}
```

```
trait Future {  
    type Item;  
  
    fn schedule (&mut self,  
                f: Box<FnOnce (T) >) ;  
}
```

Servers as futures

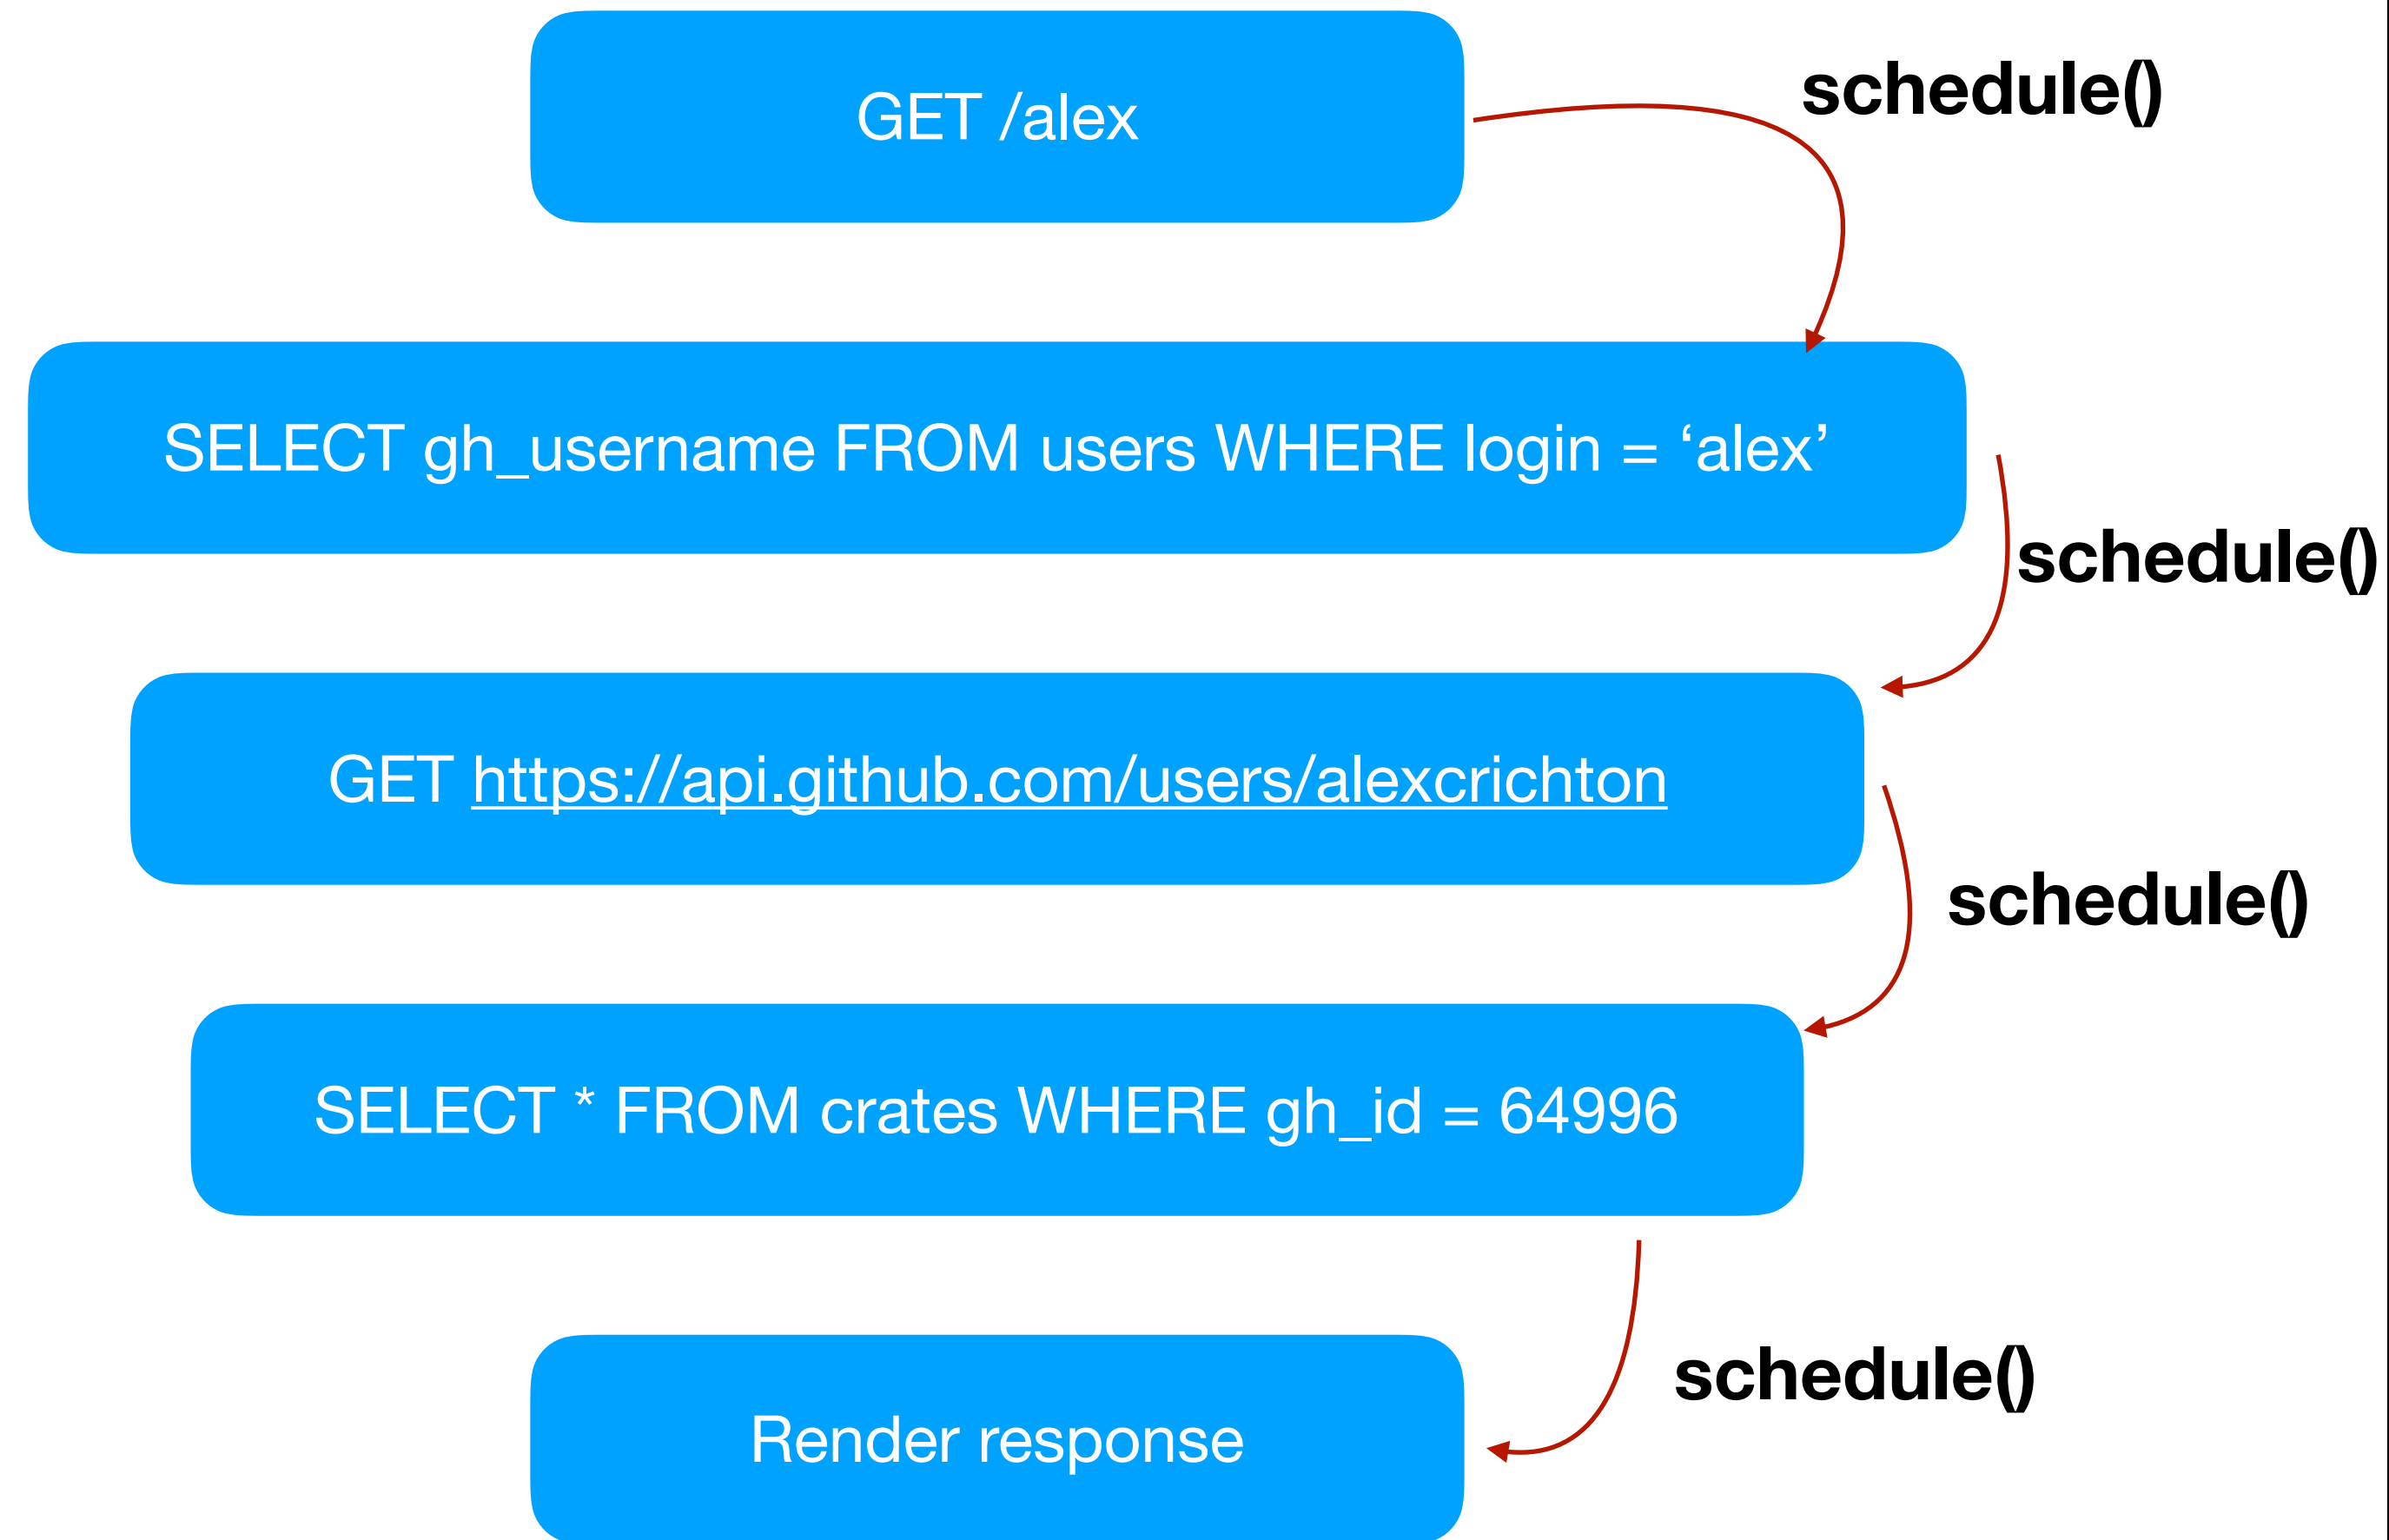
```
fn process(req: Request)
    -> impl Future<Response>
{
    // ...
}
```

Servers as futures

- Receive a request
 - Load user from the database
 - Query user's GitHub username
 - Look up username in database
 - Render a response
- 
- The word "Futures" is positioned to the right of the list, with three arrows pointing from it to the tasks "Load user from the database", "Query user's GitHub username", and "Look up username in database".

Servers as futures

impl Future<Response>



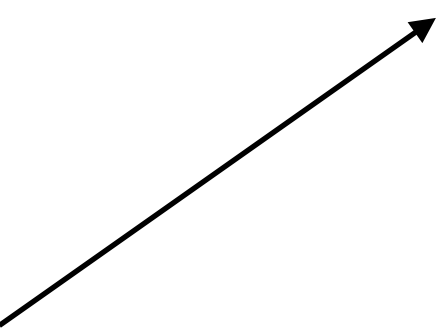
Servers as futures

- Server responses may want to be entirely a future
- Futures are frequently composed of many futures
- Futures internally make many state transitions
- **Each state transition is allocating a callback**

Futures and threading

```
trait Future {  
    type Item;  
  
    fn schedule (&mut self,  
                f: Box<FnOnce (T) >) ;  
}
```

Send? Sync?



Drawbacks of callbacks

- Expensive at runtime, forced allocation per state transition
- Gymnastics with threading
 - synchronization may be required where not necessary
 - may require multiple “threadsafe future” and “non-threadsafe future” traits

Constraints so far

- Futures must be a trait
- The Future trait must allow virtual dispatch
- State transitions need to be cheap
- Threadsafe futures shouldn't require gymnastics

impl Future<Response>

GET /alex

SELECT gh_username FROM users WHERE login = 'alex'

GET <https://api.github.com/users/alexcrichton>

SELECT * FROM crates WHERE gh_id = 64996

Render response

```
graph TD; A[GET /alex] --> B[SELECT gh_username FROM users WHERE login = 'alex']; B --> C[GET https://api.github.com/users/alexcrichton]; C --> D[SELECT * FROM crates WHERE gh_id = 64996]; D --> E[Render response];
```

Task

Read TCP connection
Future<Response>

Decrypt SSL

Decode HTTP

impl Future<Response>

Encode HTTP

Encrypt SSL

Write TCP connection
Future<Response>

One Server

Task

Task

Task

Task

Task

Task

Futures and Tasks

- A *Task* is composed of many futures
- All futures internal to a *task* track the same task
- Tasks are a unit of concurrency
- Very similar to green/lightweight threads!

```
trait Future {  
    type Item;  
  
    // ...  
}
```

“A future is a sentinel for something being computed”

```
trait Future {  
    type Item;  
  
    fn poll(&mut self)  
        -> Option<Self::Item>;  
}
```

Poll-based futures

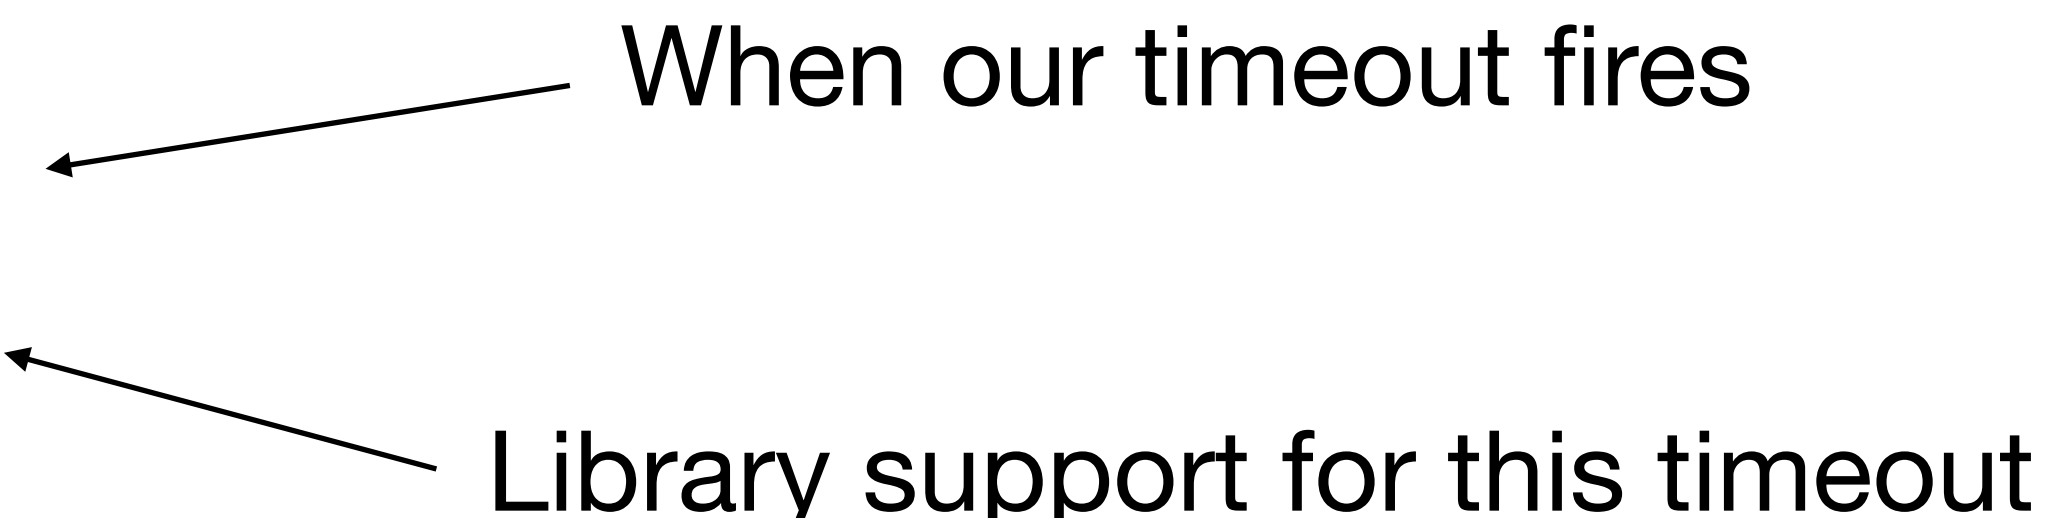
- A return value of `None` means “not ready”
- Returning `Some` resolves a future
- If you see `None`, when do you later call `poll`?

Poll-based futures

- **Futures are owned by one task**
- If a future isn't ready, the **task** needs to know when to try again
- A return value of `None` **automatically** means “I'll notify you later”

```
struct Timeout {  
    // ...  
}
```

```
struct Timeout {  
    when: Instant,  
    timer: Timer,  
}
```



When our timeout fires

Library support for this timeout

```
// provided by another library  
impl Timer {  
    fn defer(&self,  
            at: Instant,  
            f: impl FnOnce()) {  
        // ...  
    }  
}
```

```
impl Future for Timeout {  
    type Item = ();  
  
    fn poll(&mut self) -> Option<()> {  
        // ...  
    }  
}
```



```
impl Future for Timeout {  
    type Item = ();  
  
    fn poll(&mut self) -> Option<()> {  
        if self.when < Instant::now() {  
            return Some(())  
        }  
        // ...  
    }  
}
```

```
use futures::task;

impl Future for Timeout {
    type Item = ();

    fn poll(&mut self) -> Option<()> {
        if self.when < Instant::now() {
            return Some(())
        }
        let task = task::current();
        self.timer.defer(self.when, || {
            task.notify();
        });
        None // not ready yet
    }
}
```

Constraints



`trait Future`



Virtual dispatch

- Cheap
- Threadsafe

```
trait Future {  
    type Item;  
  
    fn poll(&mut self)  
        -> Option<Self::Item>;  
}
```

Constraints

✓ `trait Future`

✓ Virtual dispatch

✓ Cheap

- Threadsafe

```
mod task {  
    fn current() -> Task {  
        // just an Arc bump  
    }  
  
    impl Task {  
        fn notify(&self) {  
            // just a queue  
        }  
    }  
}
```

Constraints

✓ `trait Future`

✓ Virtual dispatch

✓ Cheap

✓ Threadsafe

```
fn assert<T: Send + Sync>() {}
```

```
assert::<Task>();
```

I/O and Polling futures

- In Tokio **everything** is a future
- Futures may eventually need to wait for I/O
- Tokio's job is to route I/O notifications to tasks

I/O and Polling futures

```
// async  
let amt = socket.read(buffer);  
assert_eq!(amt, EWOULDBLOCK);
```

- I/O all returns immediately
- “would block” turns into notifying a task
- “would block” translated to `NotReady`

I/O and Polling futures

```
trait AsyncRead: Read {  
    // empty!  
}  
  
trait AsyncWrite: Write {  
    fn shutdown(&mut self)  
        -> Poll<(), io::Error>;  
}
```

Marker traits for the current future's task is scheduled to receive a notification when the object is ready again

I/O and Polling futures

```
impl<E: Read> Read for PollEvented<E>{  
    fn read(&mut self, buf: &mut [u8])  
        -> io::Result<usize>  
    {  
        if self.poll_read().is_not_ready() {  
            return Err(wouldblock())  
        }  
        let r = self.get_mut().read(buf);  
        if is_wouldblock(&r) {  
            self.need_read();  
        }  
        return r  
    }  
}
```

Tokio's Event Loop

Kernel: file descriptor 5 is ready for reading

Tokio: Ok, let me go wake up that task.

Tokio's Event Loop

- Responsible for blocking the current thread
- Dispatches events received from the kernel
- Translates I/O notifications to task notifications

A man with short brown hair, wearing a brown leather jacket over a patterned shirt and blue jeans, stands in front of a futuristic car. He is holding a pair of sunglasses on his head with his right hand. The car has a large, glowing orange light on its side. The background is a dark, blue, and purple sky with a bright orange light source on the right.

doc.rust-lang.org/stable/book

users.rust-lang.org

github.com/alexcrichton/futures-rs